

Neural networks I: Perceptron

Semester 1, 2021

Kris Ehinger

Outline

- Representation learning
- Introduction to neural networks
- Perceptron
- Multilayer perceptron

Representation learning

“yes”



“no”



Example: Text representation

- How to represent the main idea of a text?
- Simple option: “bag of words”
 - Vector representing word frequencies
 - Values within the vector represent word count
 - Discards word order and context

[illegible]

[illegible]


Example: Text representation

- “Bag of words” works quite well for many tasks
- Problems?
 - Missing context, order, understanding of synonyms and phrases
 - Curse of dimensionality -- “words” space is very high-dimensional (~170,000 words in English)


Example: Image representation

- How to represent the main idea of an image?
- Images are made up of pixel values, each pixel has an RGB value
- How much information can you get from “bags of pixels”?

Pixels



A



B

Example: Image representation

- Pixel-level representations can work for constrained tasks
- But more complex visual tasks require more complex features (shapes, objects)
 - How to define these?
- Curse of dimensionality – raw pixel space is very high-dimensional (~700,000 dimensions in the example images)

Representation learning in NNs

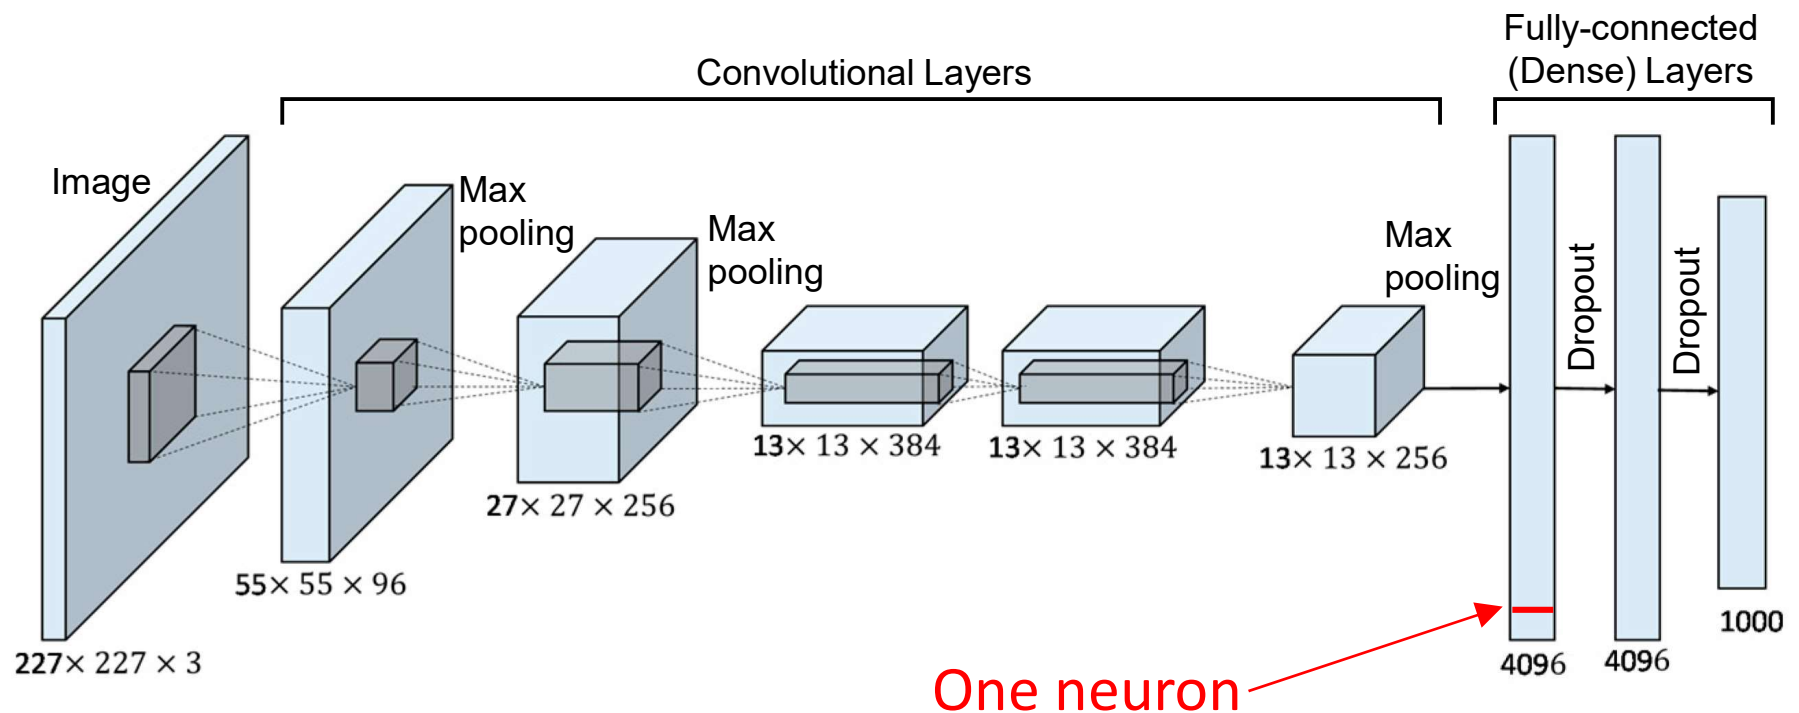
- Representation learning is a common application for neural networks
- Networks learn a feature hierarchy – from simple combinations of the input to more complex features (sometimes called **embeddings**)
- Embeddings are:
 - Low-dimensional representations of the input
 - Often useful for a range of tasks (not just the task on which the network was originally trained)



Introduction to neural networks

Convolutional neural network

“AlexNet”: Krizhevsky, Sutskever, & Hinton (2012)



Biological basis

- Hebbian learning (Hebb, 1949) – model for how neural connections change during learning
- “neurons wire together if they fire together” (Löwel & Singer, 1992)
- Over time, more weight on features associated with a target (like a class label), low weight on features not associated with the target

Biological neuron





Image: *Anatomy and Physiology*, US National Cancer Institute SEER Program

Artificial neuron

Inputs



$$y_i = f \left(\left[\sum_j w_j x_{ij} \right] + b \right) = f(\mathbf{w} \cdot \mathbf{x}_i + b)$$

Artificial neuron

- The basic unit of a neural network is the neuron, which is defined as follows:
 - input = a vector \mathbf{x}_i ($\langle x_{i1}, x_{i2}, \dots, x_{in} \rangle \in \mathbb{R}^n$)
 - output = a scalar $y_i \in \mathbb{R}$
 - hyper-parameter: an activation function f
 - parameters: a vector of weights \mathbf{w} ($\langle w_1, w_2, \dots, w_n \rangle \in \mathbb{R}^n$) plus a bias term b ($b \equiv w_0$)
- Mathematically:

$$y_i = f\left(\left[\sum_j w_j x_{ij}\right] + b\right) = f(\mathbf{w} \cdot \mathbf{x}_i + b)$$




Perceptron

Perceptron

- Perceptron = a neural network with just a single neuron
- A perceptron is a binary linear classifier, which can be written as

$$f(\mathbf{w} \cdot \mathbf{x}_i + b) = \begin{cases} 1 & \text{if } \mathbf{w} \cdot \mathbf{x}_i + b \geq 0 \\ 0 & \text{otherwise} \end{cases}$$



Cornell University News Service records, #4-3-15. Division of Rare and Manuscript Collections, Cornell University Library.
<https://digital.library.cornell.edu/catalog/ss:550351>

Training a perceptron

- Training a perceptron entails finding the weights \mathbf{w} which minimize errors on the training data
- The classic way to train a perceptron is by iterating over examples in a training dataset. Each iteration over the whole dataset is called an **epoch**.
- Perceptron algorithm
 - For training examples, compute y based on current weights
 - Update weights based on the difference between predicted y and true label

Training a perceptron

Initialize weight vector \mathbf{w} to random values

repeat

for each training instance (\mathbf{x}_i, y_i) **do**

 compute $\hat{y}_i = f(\mathbf{w} \cdot \mathbf{x}_i + b)$

for each weight w_j **do**

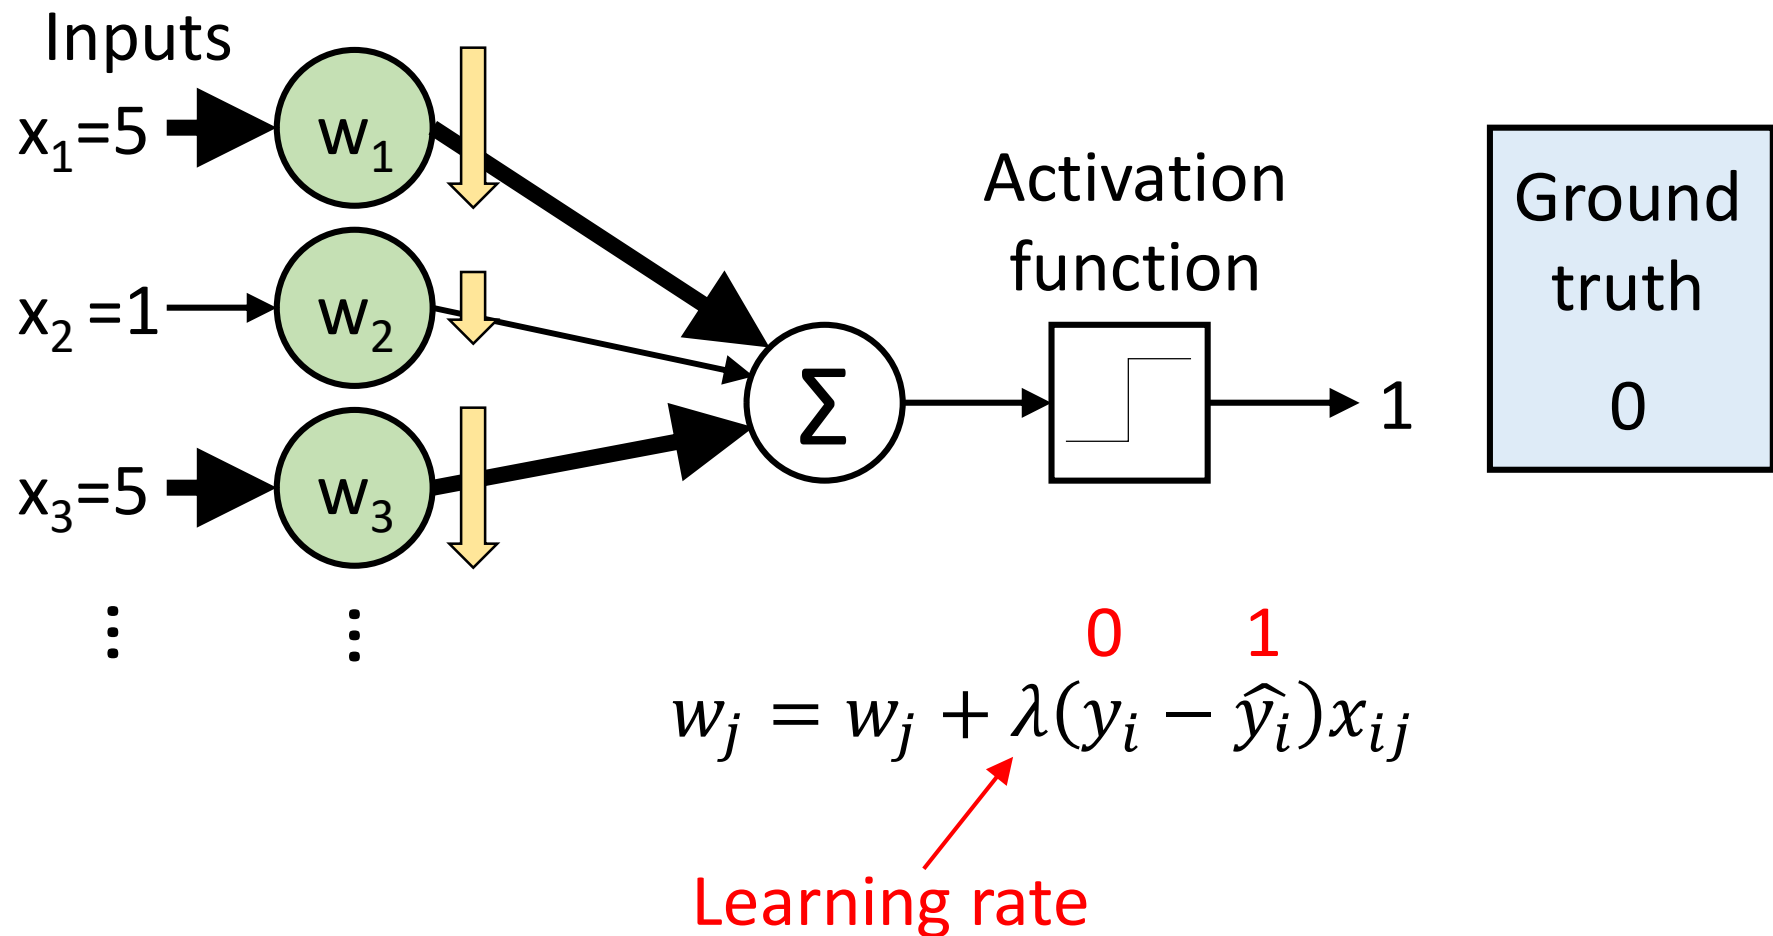
 update $w_j \leftarrow w_j + \underline{\lambda(y_i - \hat{y}_i)x_{ij}}$

end for


end for

until stopping condition is met

Weight update



Weight update



$$w_j = w_j + \overset{0}{\lambda}(y_i - \overset{1}{\hat{y}_i})x_{ij}$$

Learning rate

Numeric example

Training instances

$\langle x_{i1}, x_{i2} \rangle$	y_i
1,1	1
1,2	1
0,0	0
-1,0	0

$\langle b, w_1, w_2 \rangle = 0, 0, 0$

Epoch 1:

$\langle x_{i1}, x_{i2} \rangle$	$b + w_1 \cdot x_{i1} + w_2 \cdot x_{i2}$	\hat{y}_i	y_i
1,1	$0 + 0 + 0 = 0$	1	1

Learning rate

$\lambda = 1$

$$f(\mathbf{w} \cdot \mathbf{x}_i + b) = \begin{cases} 1 & \text{if } \mathbf{w} \cdot \mathbf{x}_i + b \geq 0 \\ 0 & \text{otherwise} \end{cases}$$

Numeric example

Training instances

$\langle x_{i1}, x_{i2} \rangle$	y_i
1,1	1
1,2	1
0,0	0
-1,0	0

$\langle b, w_1, w_2 \rangle = -1, 0, 0$

Epoch 2:

$\langle x_{i1}, x_{i2} \rangle$	$b + w_1 \cdot x_{i1} + w_2 \cdot x_{i2}$	\hat{y}_i	y_i
1,1	$-1 + 0 + 0 = -1$	0	1

Learning rate

$\lambda = 1$

$$f(\mathbf{w} \cdot \mathbf{x}_i + b) = \begin{cases} 1 & \text{if } \mathbf{w} \cdot \mathbf{x}_i + b \geq 0 \\ 0 & \text{otherwise} \end{cases}$$

Numeric example

Training instances

$\langle x_{i1}, x_{i2} \rangle$	y_i
1,1	1
1,2	1
0,0	0
-1,0	0

$\langle b, w_1, w_2 \rangle = -1, 1, 1$

Epoch 3:

$\langle x_{i1}, x_{i2} \rangle$	$b + w_1 \cdot x_{i1} + w_2 \cdot x_{i2}$	\hat{y}_i	y_i
1,1	$-1 + 1 + 1 = 1$	1	1

Learning rate

$\lambda = 1$

Convergence – there were no updates
in this epoch, so end training

$$f(\mathbf{w} \cdot \mathbf{x}_i + b) = \begin{cases} 1 & \text{if } \mathbf{w} \cdot \mathbf{x}_i + b \geq 0 \\ 0 & \text{otherwise} \end{cases}$$


Activation function

- What does it do?
- In a perceptron
 - Maps the linear response to the range we want (class labels 0 or 1)
- In a multilayer perceptron / neural network:
 - Adds a non-linearity after each linear operation
 - More about this later!

Activation function


- Common choices:
- (logistic) sigmoid (σ):

$$f(x) = \frac{1}{1 + e^{-x}}$$




- hyperbolic tan (tanh):

$$f(x) = \frac{e^{2x} - 1}{e^{2x} + 1}$$



- rectified linear unit (ReLU):

$$f(x) = \max(0, x)$$



Perceptron properties

- The perceptron algorithm is guaranteed to converge for linearly-separable data, but the convergence point (class boundary) will depend on:
 - The initial values of the weights and bias
 - The learning rate
- Not guaranteed to maximise the margin between classes
- Not guaranteed to converge over non-linearly-separable data

Perceptron properties

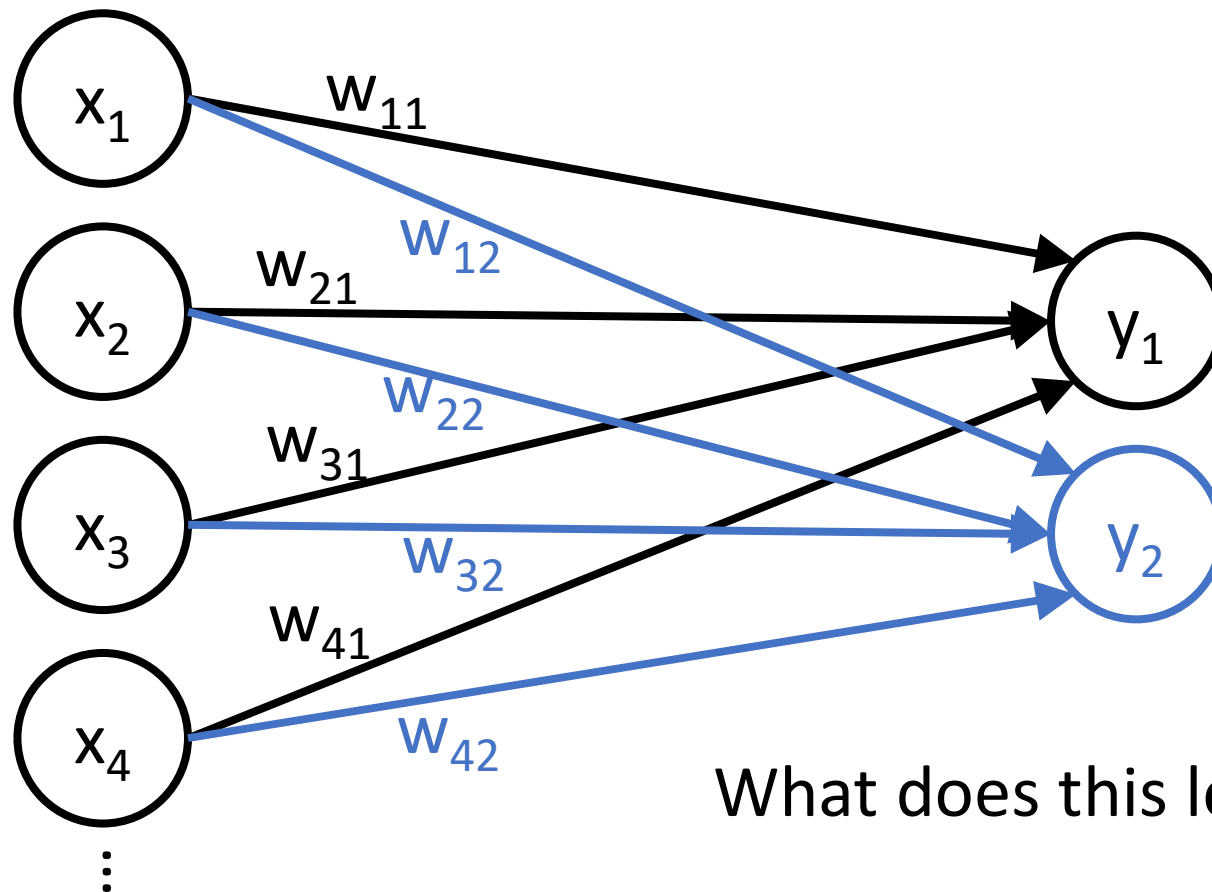
- A perceptron with a sigmoid activation function (plus a binary step to convert the output to 0 or 1) is equivalent to logistic regression:

$$y_i = \frac{1}{1 + e^{-(\mathbf{w} \cdot \mathbf{x}_i + b)}}$$

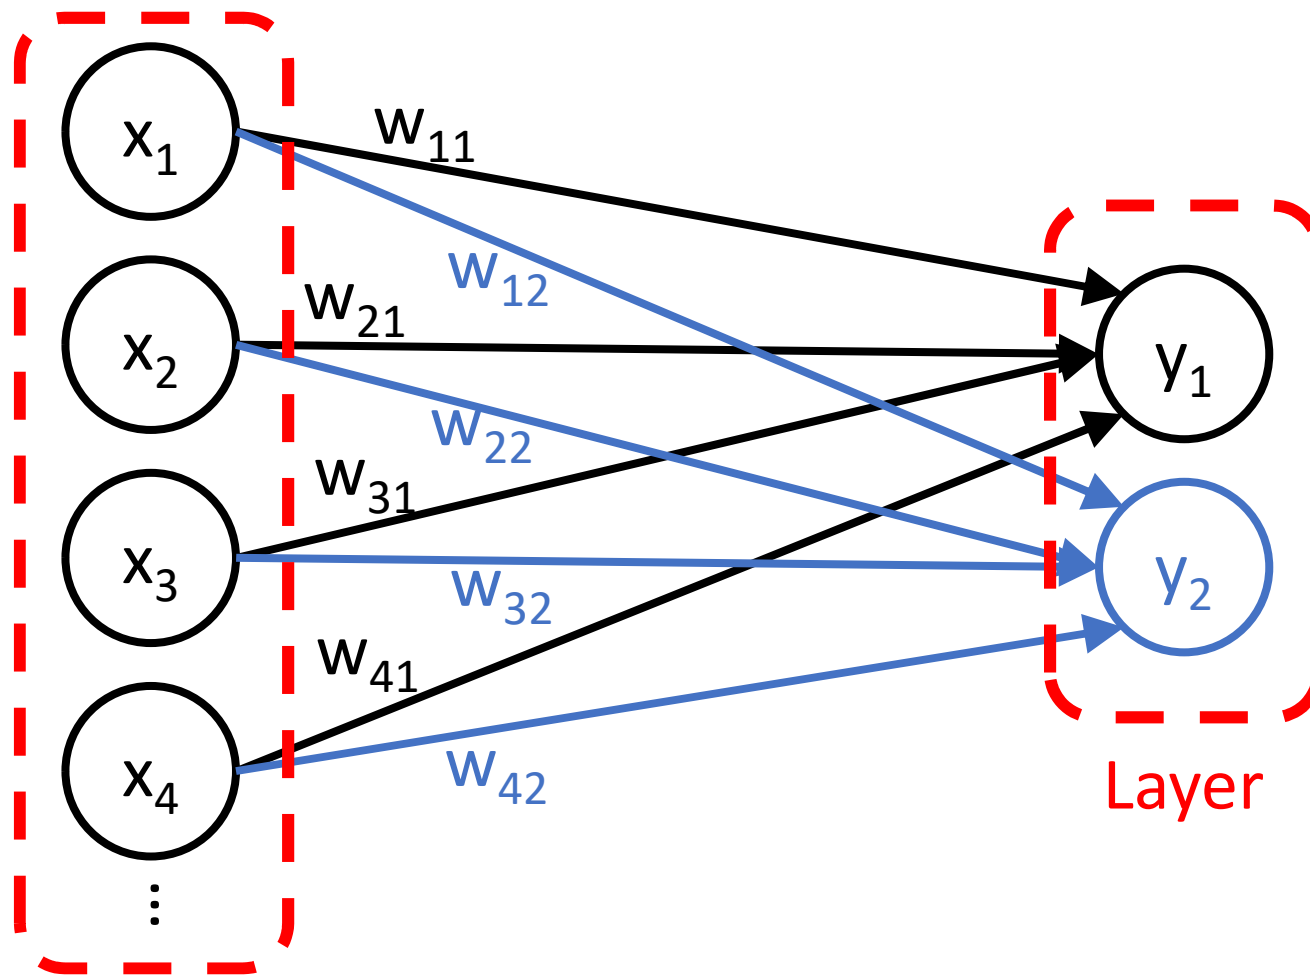


Multilayer perceptron


Extension 1: Multiple outputs



Extension 1: Multiple outputs



Extension 2: Multiple layers




What does this learn?

Multilayer perceptron

- Layer = a group of neurons working in parallel on the same input
- Multilayer perception (MLP):
 - An **input layer**
 - At least one **hidden layer**. Each hidden layer receives the previous layer's output as its input
 - An **output layer** which receives input from a hidden layer and outputs the class label

Hidden layer

- What does it learn?
 - Representation learning – hidden layers learn some features that are useful for the output layer



https://scikit-learn.org/stable/auto_examples/neural_networks/plot_mnist_filters.html


Hidden layer

- How many neurons?
 - Given that we don't know exactly what features this layer should learn, how to pick the right number of neurons?
 - In theory, depends on factors like complexity of decision boundary
 - In practice, often just pick an arbitrary value in between the input and output size

Hidden layer

- How does it learn?
 - There's not just one “correct” output for the neurons in a hidden layer, like there is for the output layer
 - How can the hidden layer(s) be trained?
- Solution: backpropagation
 - **Stochastic** gradient descent
 - Find parameters (weights+biases) to minimize loss on the training dataset

Gradient descent



Backpropagation

- Form of gradient descent – iterative process for finding the weight parameters that minimize error
- Just like single perceptron training, requires a learning rate (hyperparameter)
- Idea: look at the derivative of error and update weights in the direction that would reduce error
 - Compute errors at the output layer w.r.t. each weight using partial differentiation
 - Propagate errors back to each of the earlier layers
 - Chain rule used to efficiently compute derivatives

Output layer

- Structure depends on task
- Common options:
 - Binary classification – 1 output neuron with step activation function
 - N-way classification – N output neurons and softmax activation function
 - Regression – 1 output neuron with identity activation function

Neural network/MLP properties

- Universal approximation theorem: a feed-forward neural network with a single hidden layer (with finite neurons) is able to approximate any continuous function on \mathbb{R}^n
- Means the network can learn any (linear or non-linear) basis function dynamically, unlike many other methods (e.g., SVM where kernel is a hyperparameter)
- Note that this requires non-linearities

Non-linearities

- With a non-linear activation function:

$$f(\mathbf{w}_{L2} \cdot f(\mathbf{w}_{L1} \cdot f(\mathbf{w}_{L0} \cdot \mathbf{x}_i + \mathbf{b}_{L0}) + \mathbf{b}_{L1}) + \mathbf{b}_{L2})$$

- Without non-linearity:

$$\begin{aligned} \mathbf{w}_{L2} \cdot (\mathbf{w}_{L1} \cdot (\mathbf{w}_{L0} \cdot \mathbf{x}_i + \mathbf{b}_{L0}) + \mathbf{b}_{L1}) + \mathbf{b}_{L2} \\ = \mathbf{w}_{\alpha} \cdot \mathbf{x}_i + \mathbf{b}_{\alpha} \end{aligned}$$

- Multiple layers don't accomplish anything – the same computations could have been done in one layer

Neural network/ MLP properties

- True or false?
- Since NNs are universal approximators, they are guaranteed to generalise better than any other machine learning method.
- Since NNs learn their own representations in hidden layers, they don't require any feature engineering.

Neural network/ MLP properties

- MLP advantages
 - Can be adapted to many types of problems (classification, regression)
 - Universal approximator – can model arbitrary basis functions
 - Representation learning in hidden layers
- MLP disadvantages
 - Very high number of parameters – slow to train, prone to overfitting, high memory requirements
 - Stochastic gradient descent – not guaranteed to converge to the same solution every time